

BLG 431E SYSTEM PROGRAMMING

PROJECT 1

Group 57

Ahmet Göktuğ SEVİNÇ – 150140120

Abdullah Melih Canal – 150130056

In the project we were asked to implement a new system call (*set_myFlag*) to set the value of a flag (*myFlag*) which is added into task descriptor of a process. To be able to add a new flag into the processes we had to change the task descriptor of the processes. In Linux, the kernel stores the list of processes in a circular doubly link list called the *task list*. Each element in the task list is a *process descriptor* which contains all of the information about a process and these process descriptors are the type of *struct task_struct*. Therefore, firstly, we added a new field into that *task_struct* structure which is defined in `<linux/sched.h>` and that header file was located in `/include/linux/sched.h` path :

```
1053 struct task_struct {
1054     volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
```

...

```
1464 #if defined(CONFIG_BCACHE) || defined(CONFIG_BCACHE_MODULE)
1465     unsigned int sequential_io;
1466     unsigned int sequential_io_avg;
1467 #endif
1468
1469     int myFlag;
1470 };
```

After adding our flag to the task descriptor, to initialize our new field during system initialization (creation of process 0) we needed to update `INIT_TASK` macro located in `/include/linux/init_task.h` :

```
162 #define INIT_TASK(tsk) \
163 { \
164     .state = 0, \
165     .stack = &init_thread_info, \
```

...

```
217     thread_node = LIST_HEAD_INIT(init_signals.thread_head), \
218     .myFlag = 0, \
219     INIT_IDS \
220     INIT_PERF_EVENTS(tsk) \
```

After those initial operations, we wrote our system call function and added it to the kernel. For this purpose, we created a new folder (*set_myFlag*) under root of the source and in that folder we created our new system call file (*set_myFlag.c*) :

```

1  #include <linux/syscalls.h>
2  #include <linux/kernel.h>
3  #include <asm/errno.h>
4
5  asmlinkage long sys_set_myFlag(pid_t pid, int flag_val)
6  {
7      if (capable(CAP_SYS_ADMIN))
8      {
9          struct task_struct *tsk = find_task_by_vpid(pid);
10
11         if (tsk != NULL)
12         {
13             if (flag_val == 1 || flag_val == 0)
14             {
15                 tsk->myFlag = flag_val;
16                 return 0;
17             }
18             return -EINVAL;
19         }
20         return -ESRCH;
21     }
22     return -EPERM;
23 }

```

First of this system call can only be used by processes with root privileges so first *if* statement ensures that restriction. If the calling process does not have the root privileges we returned an error (*-EPERM*) to indicate that it is not permitted to do that operation. Then to be able to obtain a pointer to the task descriptor of a process with its pid we examined *sched.h* file and found the appropriate function *find_task_by_vpid()* which takes a pid as a parameter. Then we simply checked whether we obtained the pointer or not. If not, we returned an error (*-ESRCH*) to indicate a process with that pid could not be found. And finally, we checked the given flag value to the function. If it is not zero (0) or one (1), we returned another error (*-EINVAL*) to indicate value is invalid. If all constraints are satisfied, we assigned flag value to the current process.

After writing our system call, we also created a *Makefile* under that folder and added it to the Makefile of the source to make this call available at the compilation.

/set_myFlag/Makefile:

```

1  obj-y := set_myFlag.o

```

/Makefile:

```

537  drivers-y := drivers/ sound/ firmware/ ubuntu/
538  net-y      := net/
539  libs-y     := lib/
540  core-y     := usr/ set_myFlag/

```

Then we modified system call table under */arch/x86/syscalls/syscall_32.tbl* and system call header file under */include/linux/syscall.h*.

/arch/x86/syscalls/syscall_32.tbl:

```
362 353 i386 renameat2 sys_ni_syscall
363 354 i386 seccomp sys_seccomp
364 355 i386 set_myFlag sys_set_myFlag
365
```

/include/linux/syscall.h:

```
850 asmlinkage long sys_seccomp(unsigned int op, unsigned int flags,
851 const char __user *uargs);
852 asmlinkage long sys_set_myFlag(pid_t pid, int flag);
853 #endif
```

Now, since all of the operations regarding system call was finished, we could start to second requirement of the project: modifying fork and exit system calls. In Linux, new processes are created by *fork* system call and that call uses *do_fork* function located in */kernel/fork.c*. So, we were asked to change that function so that, if the value of our new flag (*myFlag*) is zero (0), default actions will be performed but if it is one (1) and its nice value is greater than 10 no processes will be created.

```
1642 long do_fork(unsigned long clone_flags,
1643 unsigned long stack_start,
1644 unsigned long stack_size,
1645 int __user *parent_tidptr,
1646 int __user *child_tidptr)
1647 {
1648     struct task_struct *tsk = current;
1649     // Do not create a child process under this situations
1650     if (tsk->myFlag == 1 && task_nice(tsk) > 10)
1651         return -ECHILD;
1652 }
```

Here, inside *do_fork* function, first of all we read the current process and checked its flag value and nice value. If its flag is 1 and nice value is greater than 10, we returned from the function with an error (*-ECHILD*) to indicate no child process will be created. If these conditions are satisfied, *do_fork* function calls another function *copy_process* to create new child process. Therefore, we simply added a line to initialize the flag of the child processes.

```
1191 static struct task_struct *copy_process(unsigned long clone_flags,
1192 unsigned long stack_start,
1193 unsigned long stack_size,
1194 int __user *child_tidptr,
1195 struct pid *pid,
1196 int trace)
1197 {
1198     int retval;
1199     struct task_struct *p;
1200 }
```

...

```
1289
1290     p->myFlag = 0;
1291 }
```

Finally, we were also asked to modify exit system call so that, if the myFlag value of process is zero (0), default actions will be performed, but if its value is one (1) and its nice value is greater than 10, not only that specific process, but all of its siblings will also be terminated. In linux processes are terminated via exit system call and for this purpose *do_exit* function is used located in */kernel/exit.c* file. So, we made some changes in that function.

```
709     if (tsk->myFlag == 1 && task_nice(tsk) > 10)
710     {
711         struct list_head *sibling_list;
712         struct task_struct * _sibling;
713         struct task_struct * _parent = current->parent;
714         list_for_each(sibling_list, &_parent->children)
715         {
716             _sibling = list_entry(sibling_list, struct task_struct, sibling);
717             siginfo_t siginfo;
718             kill_proc_info(SIGKILL, &siginfo, _sibling->pid);
719         }
720     }
```

Here, first of we checked whether the flag value of the current process is 1 and its nice value is greater than 10, or not. If these conditions were satisfied, we needed to kill all of the siblings of that process along with itself. To be able to do that, we used *list_for_each* and *list_entry* macros. Firstly we created a linked list called *sibling_list*, then we created two *task_struct* pointers to point each of siblings (*_sibling*) and the parent of the current process (*_parent*). After that we used *list_for_each* macro to iterate over all of the childrens of that parent and by using *list_entry* macro we obtained those child processes and terminated them by sending *SIGKILL* signal.

TESTING

For test operations we wrote 3 functions. Firstly, we wrote a function to set the flag (my_Flag) value of a process from the command line. *set_Flag.c* :

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <errno.h>
5  #include <string.h>
6
7  #define set_myFlag 355
8
9  int main(int argc, char* argv[]){
10     int result;
11     int pid = atoi(argv[1]);
12     int flag_value = atoi(argv[2]);
13
14     result = syscall(set_myFlag, pid, flag_value);
15     if(result == 0)
16         printf("Success\n");
17     else
18         printf("Error, %s\n", strerror(-result));
19
20     return result;
21 }
```

This function takes process id, and the flag value from the command line. Then sets the flag of the process by calling our system call and prints the result.

To test *fork()* operations we wrote another function. *Fork_Test()*:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <errno.h>
5  #include <string.h>
6  #include <signal.h>
7  #include <sys/types.h>
8
9  void mysignal(int signum){
10     printf("Received signal: %d\n", signum);
11 }
12
13 void mysigset(int num){
14     struct sigaction mysigaction;
15     mysigaction.sa_handler = (void *) mysignal;
16     mysigaction.sa_flags = 0;
17     sigaction(num, &mysigaction, NULL);
18 }
```

```

20 int main(){
21     mysigset(12);
22     int i, f=1, children[3], myOrder;
23     pause();
24
25     for (i=0; i<3; i++){
26         if(f>0)
27             f = fork();
28         if(f<0){
29             printf("fork error, %s\n", strerror(-f));
30         }
31         if(f==0)
32             break;
33         else
34             children[i] = f;
35     }

```

```

37     if(f>0){
38         printf("I am the parent process. My pid is: %d\n", getpid());
39     }
40     else if(f==0){
41         myOrder = i;
42         printf("I am the child %d, My pid is: %d My parent's pid is: %d\n", myOrder, getpid(), getppid());
43     }else{
44         printf("I am the parent process, child process could not be produced\n");
45     }
46 }

```

At the beginning of that function we call our sigset function with the signal number of 12 and then we immediately pause the process. If we send kill signal from the command line, everything will work properly and there will be 1 parent and 3 child processes and they will print their information as output. But, if we set flag of our process to 1 and set the nice value of it to above 10, no child will be produced.

```

sevinca@sevinca-VirtualBox:~/Desktop$ gcc -o fork_test Fork_Test.c
sevinca@sevinca-VirtualBox:~/Desktop$ ./fork_test

```

Here we started our function and it is paused. Waiting for a signal. Now we can check process list and find the process number of that process. To do this we used `ps -ef` command that lists current processes.

```

sevinca@sevinca-VirtualBox:~/Desktop$ ps -ef
UID          PID    PPID  C  STIME TTY          TIME CMD
root           1         0  0   00:00 ?        00:00:00 /sbin/init
root           2         0  0   00:00 ?        00:00:00 [kthreadd]
root           3         2  0   00:00 ?        00:00:00 [ksoftirqd/0]

```

...

```

sevinca  4140  3903  0 01:01 pts/6    00:00:00 ./fork_test
sevinca  4142  3885  0 01:01 pts/3    00:00:00 ps -ef
sevinca@sevinca-VirtualBox:~/Desktop$ renice 15 -p 4140
4140 (process ID) old priority 0, new priority 15
sevinca@sevinca-VirtualBox:~/Desktop$ sudo ./set 4140 1
Success
sevinca@sevinca-VirtualBox:~/Desktop$ kill -12 4140
sevinca@sevinca-VirtualBox:~/Desktop$

```

After learning our process number, we set our flag value to one and nice value of our process to 15, by using *renice* program. Then, we sent kill signal(12) to the process.

```
sevinca@sevinca-VirtualBox:~$ cd Desktop
sevinca@sevinca-VirtualBox:~/Desktop$ gcc -o fork_test Fork_Test.c
sevinca@sevinca-VirtualBox:~/Desktop$ ./fork_test
Received signal: 12
fork error, No child processes
fork error, No child processes
fork error, No child processes
Killed
sevinca@sevinca-VirtualBox:~/Desktop$
```

As expected, no child process is produced.

To test our exit system call we wrote another function. *Exit_Test.c* :

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  int main(){
6
7      int i, f=1, children[3], myOrder;
8
9      for (i=0; i<3; i++){
10         if(f>0)
11             f = fork();
12         if(f== -1){
13             printf("fork error ");
14             exit(1);
15         }
16         if(f==0)
17             break;
18         else
19             children[i] = f;
20     }
```

```
22     if(f>0){
23         int mypid = getpid();
24         printf("I am the parent process. My pid is: %d\n", mypid);
25         while(1);
26     }
27
28     else{
29         myOrder = i;
30         printf("I am the child %d, My pid is: %d My parent's pid is: %d\n", myOrder, getpid(), getppid());
31         while(1);
32     }
33 }
```

Here, we simply create 1 parent and 3 child processes and put all of them into infinite while loop.


```
sevinca@sevinca-VirtualBox:~/Desktop$ gcc -o exit_test Exit_Test.c
sevinca@sevinca-VirtualBox:~/Desktop$ ./exit_test
I am the parent process. My pid is: 4203
I am the child 2, My pid is: 4206 My parent's pid is: 4203
I am the child 1, My pid is: 4205 My parent's pid is: 4203
I am the child 0, My pid is: 4204 My parent's pid is: 4203
```

As it is seen, 1 parent and 3 child processes are created. Now since we know process numbers of our child processes, we can set one of their flag to 1 and nice value to greater than 10. But, before doing that we checked the current process list, by using `ps -ef` command again.

```
root      4129      2  0 01:00 ?        00:00:00 [kworker/0:1]
sevinca   4203   3903 23 01:05 pts/6    00:00:06 ./exit_test
sevinca   4204   4203 23 01:05 pts/6    00:00:06 ./exit_test
sevinca   4205   4203 23 01:05 pts/6    00:00:06 ./exit_test
sevinca   4206   4203 23 01:05 pts/6    00:00:06 ./exit_test
root      4207      2  0 01:05 ?        00:00:00 [kworker/0:2]
sevinca   4214   3885  0 01:05 pts/3    00:00:00 ps -ef
```

```
sevinca@sevinca-VirtualBox:~/Desktop$ renice 15 -p 4204
4204 (process ID) old priority 0, new priority 15
sevinca@sevinca-VirtualBox:~/Desktop$ sudo ./set 4204 1
Success
sevinca@sevinca-VirtualBox:~/Desktop$ kill 4204
sevinca@sevinca-VirtualBox:~/Desktop$
```

Here, we set the nice value and flag of a child process with the process number of 4204. Then we killed that child process by sending `-kill` signal to it. Then, we checked process list again.

```
root      4129      2  0 01:00 ?        00:00:00 [kworker/0:1]
sevinca   4203   3903 33 01:05 pts/6    00:00:29 ./exit_test
sevinca   4204   4203 16 01:05 pts/6    00:00:14 [exit_test] <defunct>
sevinca   4205   4203 23 01:05 pts/6    00:00:20 [exit_test] <defunct>
sevinca   4206   4203 23 01:05 pts/6    00:00:20 [exit_test] <defunct>
root      4207      2  0 01:05 ?        00:00:00 [kworker/0:2]
```

As expected, not only the process with the pid of 4204, but also its siblings (4205, 4206) are also killed. As it can be seen in above images, before using our `set_Flag` function we used `-sudo` command to get root privileges otherwise we got an error saying we don't have permission for that operation.